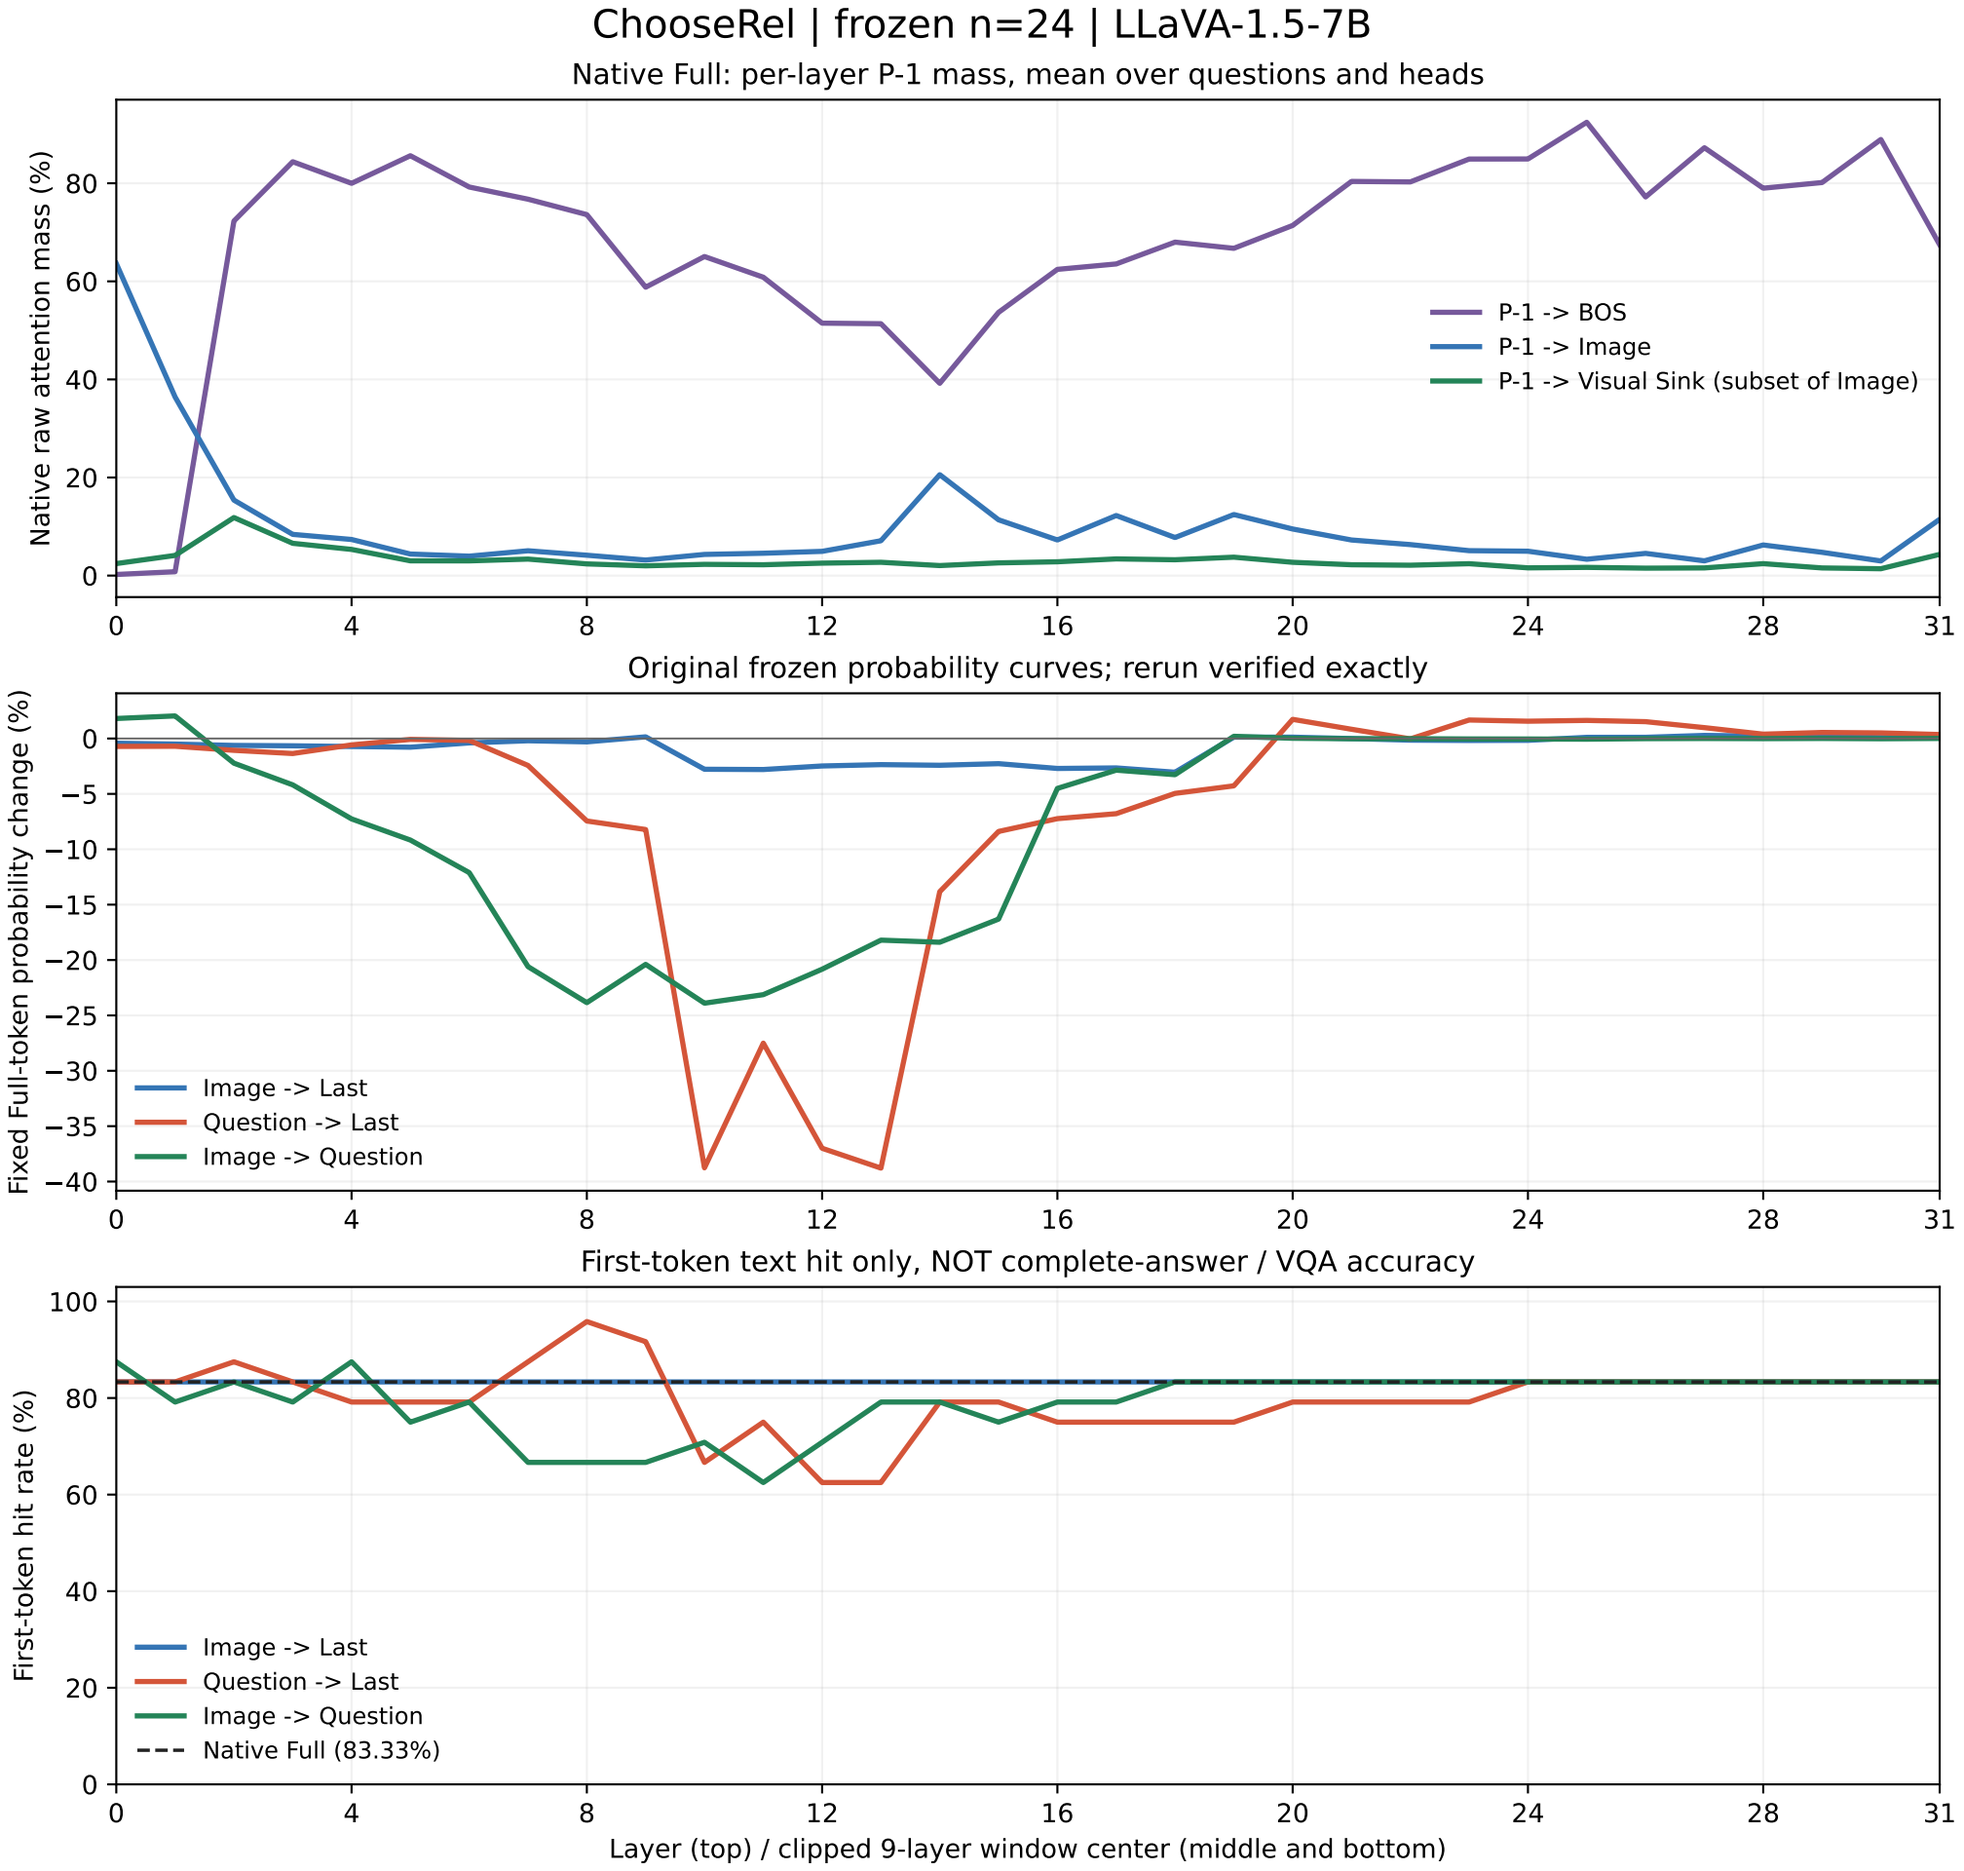
ChooseRel | frozen n=24 | LLaVA-1.5-7B
Native Full: per-layer P-1 mass, mean over questions and heads
Native raw attention mass (%)
P-1 -> BOS
P-1 -> Image
P-1 -> Visual Sink (subset of Image)
Original frozen probability curves; rerun verified exactly
Fixed Full-token probability change (%)
Image -> Last
Question -> Last
Image -> Question
First-token text hit only, NOT complete-answer / VQA accuracy
First-token hit rate (%)
Image -> Last
Question -> Last
Image -> Question
Native Full (83.33%)
Layer (top) / clipped 9-layer window center (middle and bottom)