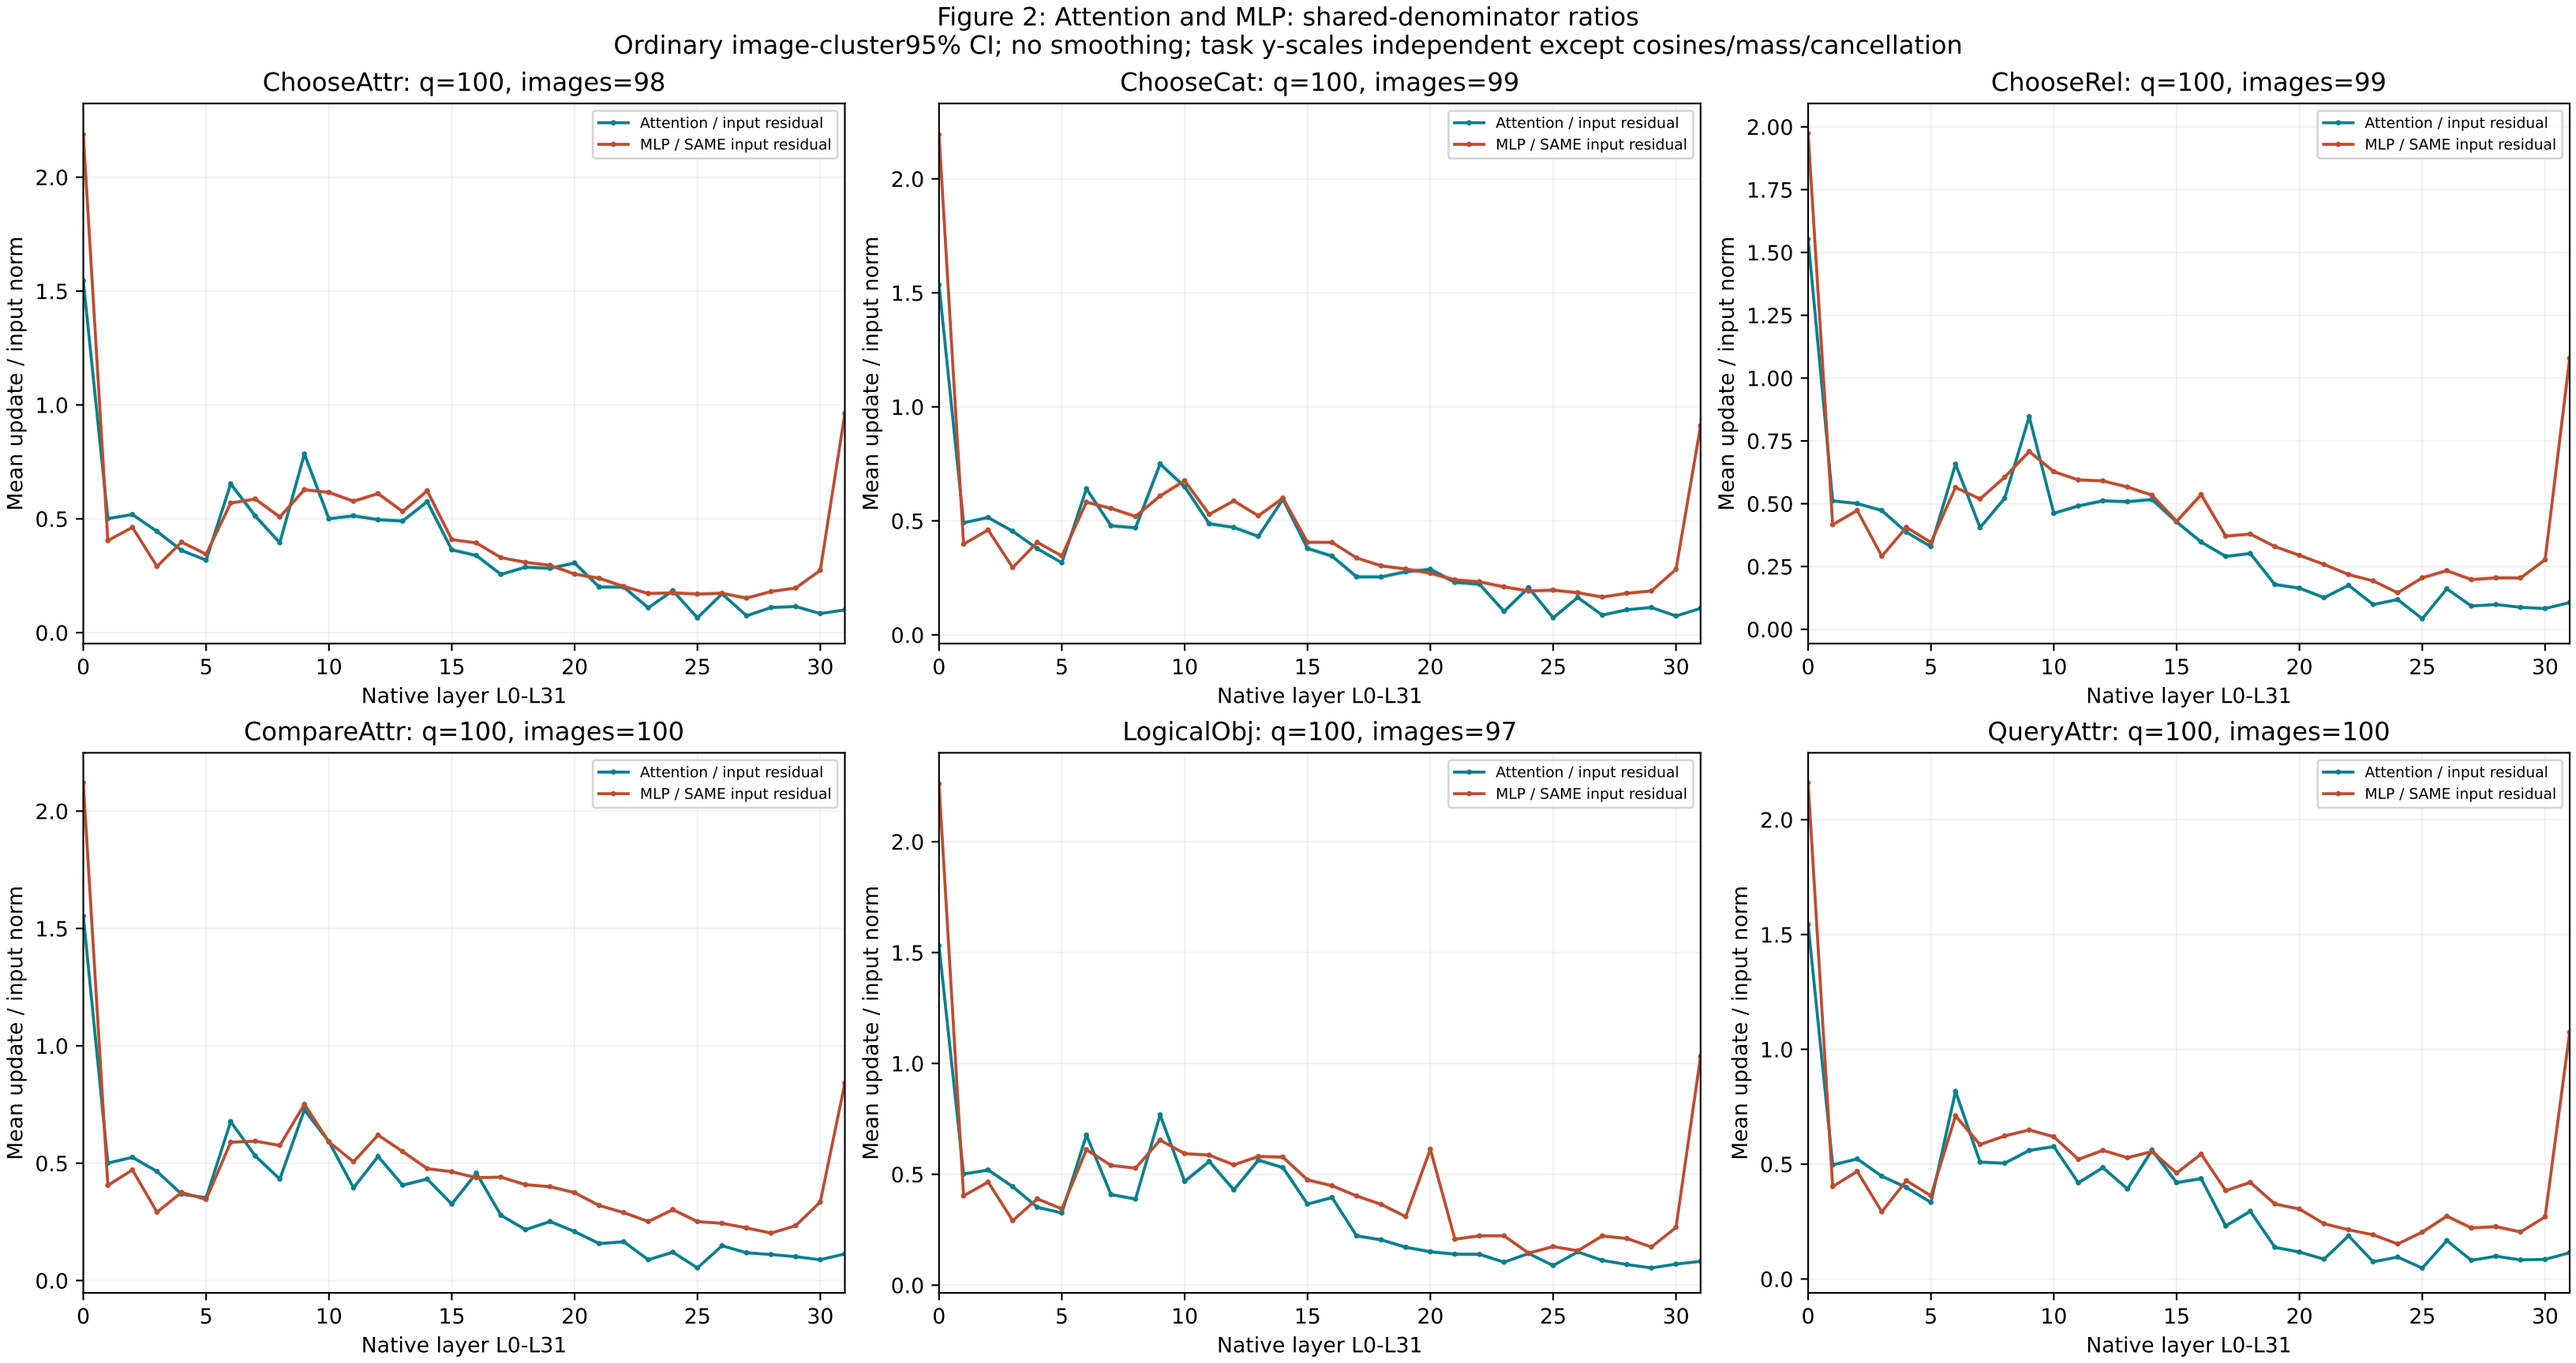
Figure 2: Attention and MLP: shared-denominator ratios
Ordinary image-cluster95% CI; no smoothing; task y-scales independent except cosines/mass/cancellation
ChooseAttr: q=100, images=98
ChooseCat: q=100, images=99
ChooseRel: q=100, images=99
CompareAttr: q=100, images=100
LogicalObj: q=100, images=97
QueryAttr: q=100, images=100
Attention / input residual
MLP / SAME input residual
Mean update / input norm
Native layer L0-L31