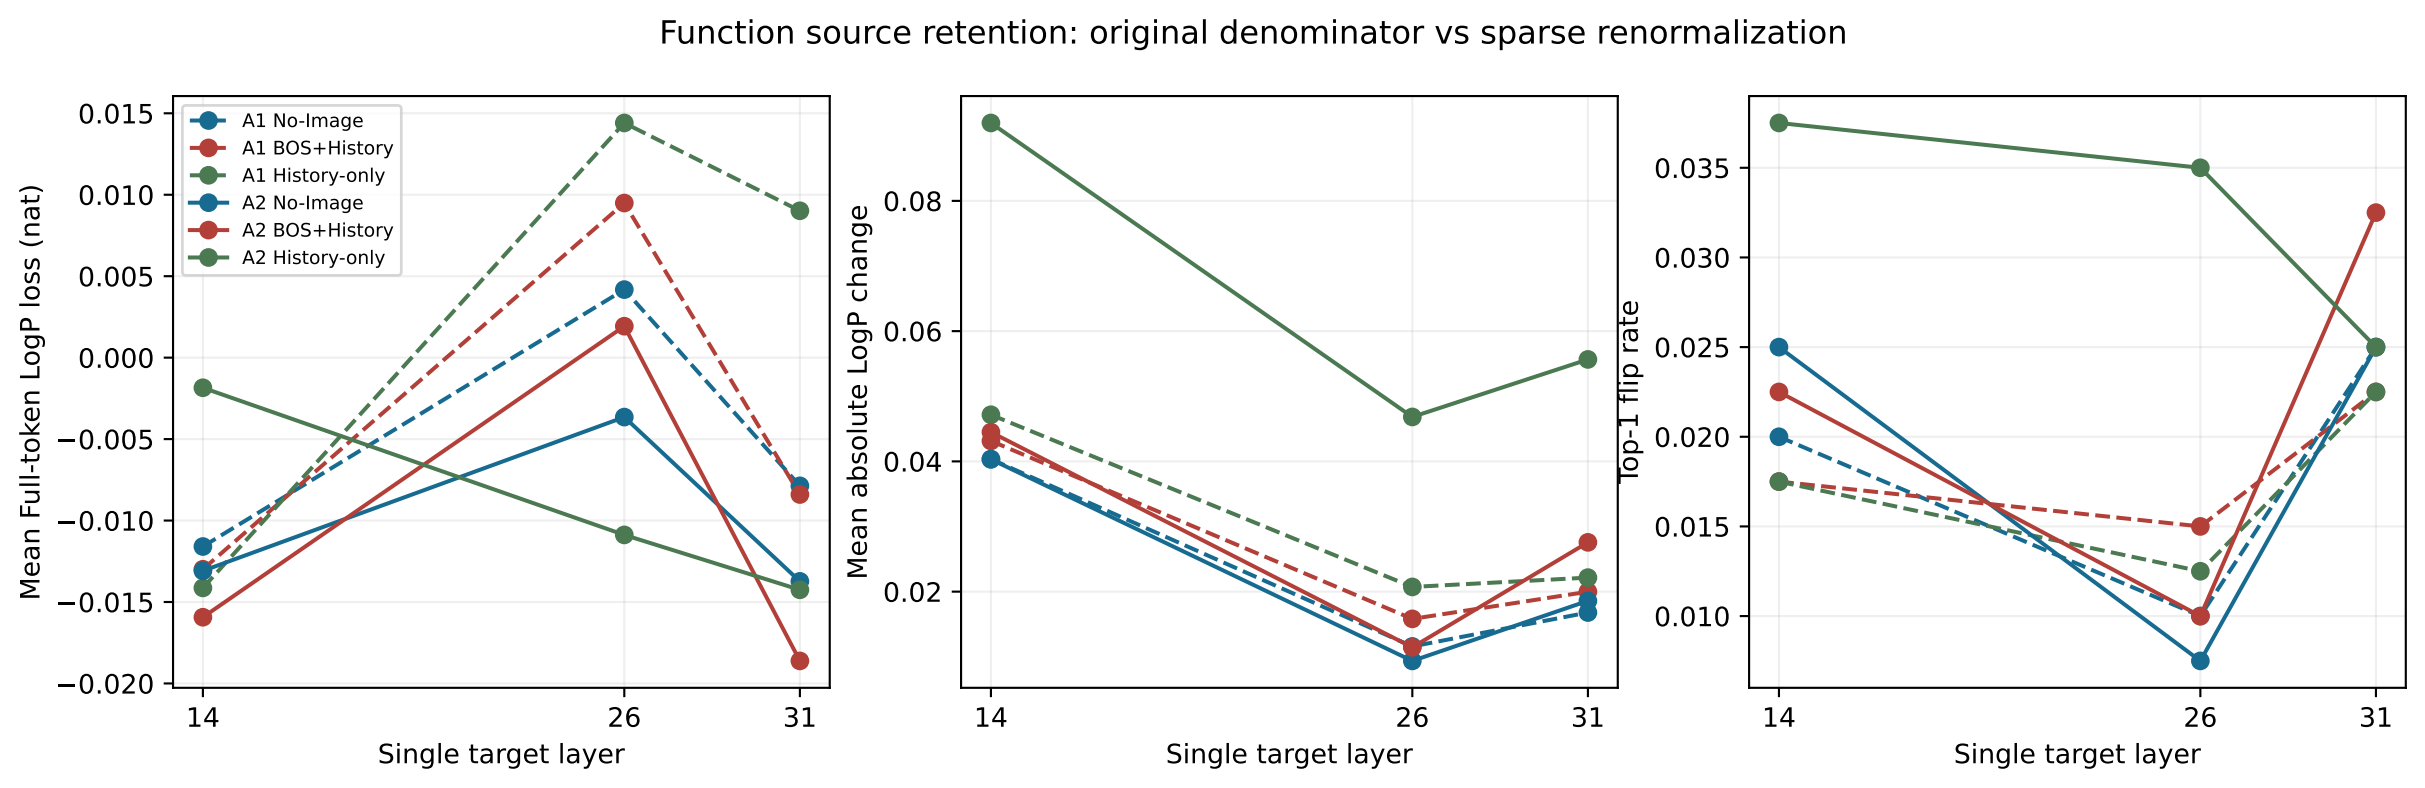
Function source retention: original denominator vs sparse renormalization
Mean Full-token LogP loss (nat)
Mean absolute LogP change
Top-1 flip rate
Single target layer
Single target layer
Single target layer
A1 No-Image
A1 BOS+History
A1 History-only
A2 No-Image
A2 BOS+History
A2 History-only
0.015
0.010
0.005
0.000
−0.005
−0.010
−0.015
−0.020
0.08
0.06
0.04
0.02
0.035
0.030
0.025
0.020
0.015
0.010
14
26
31