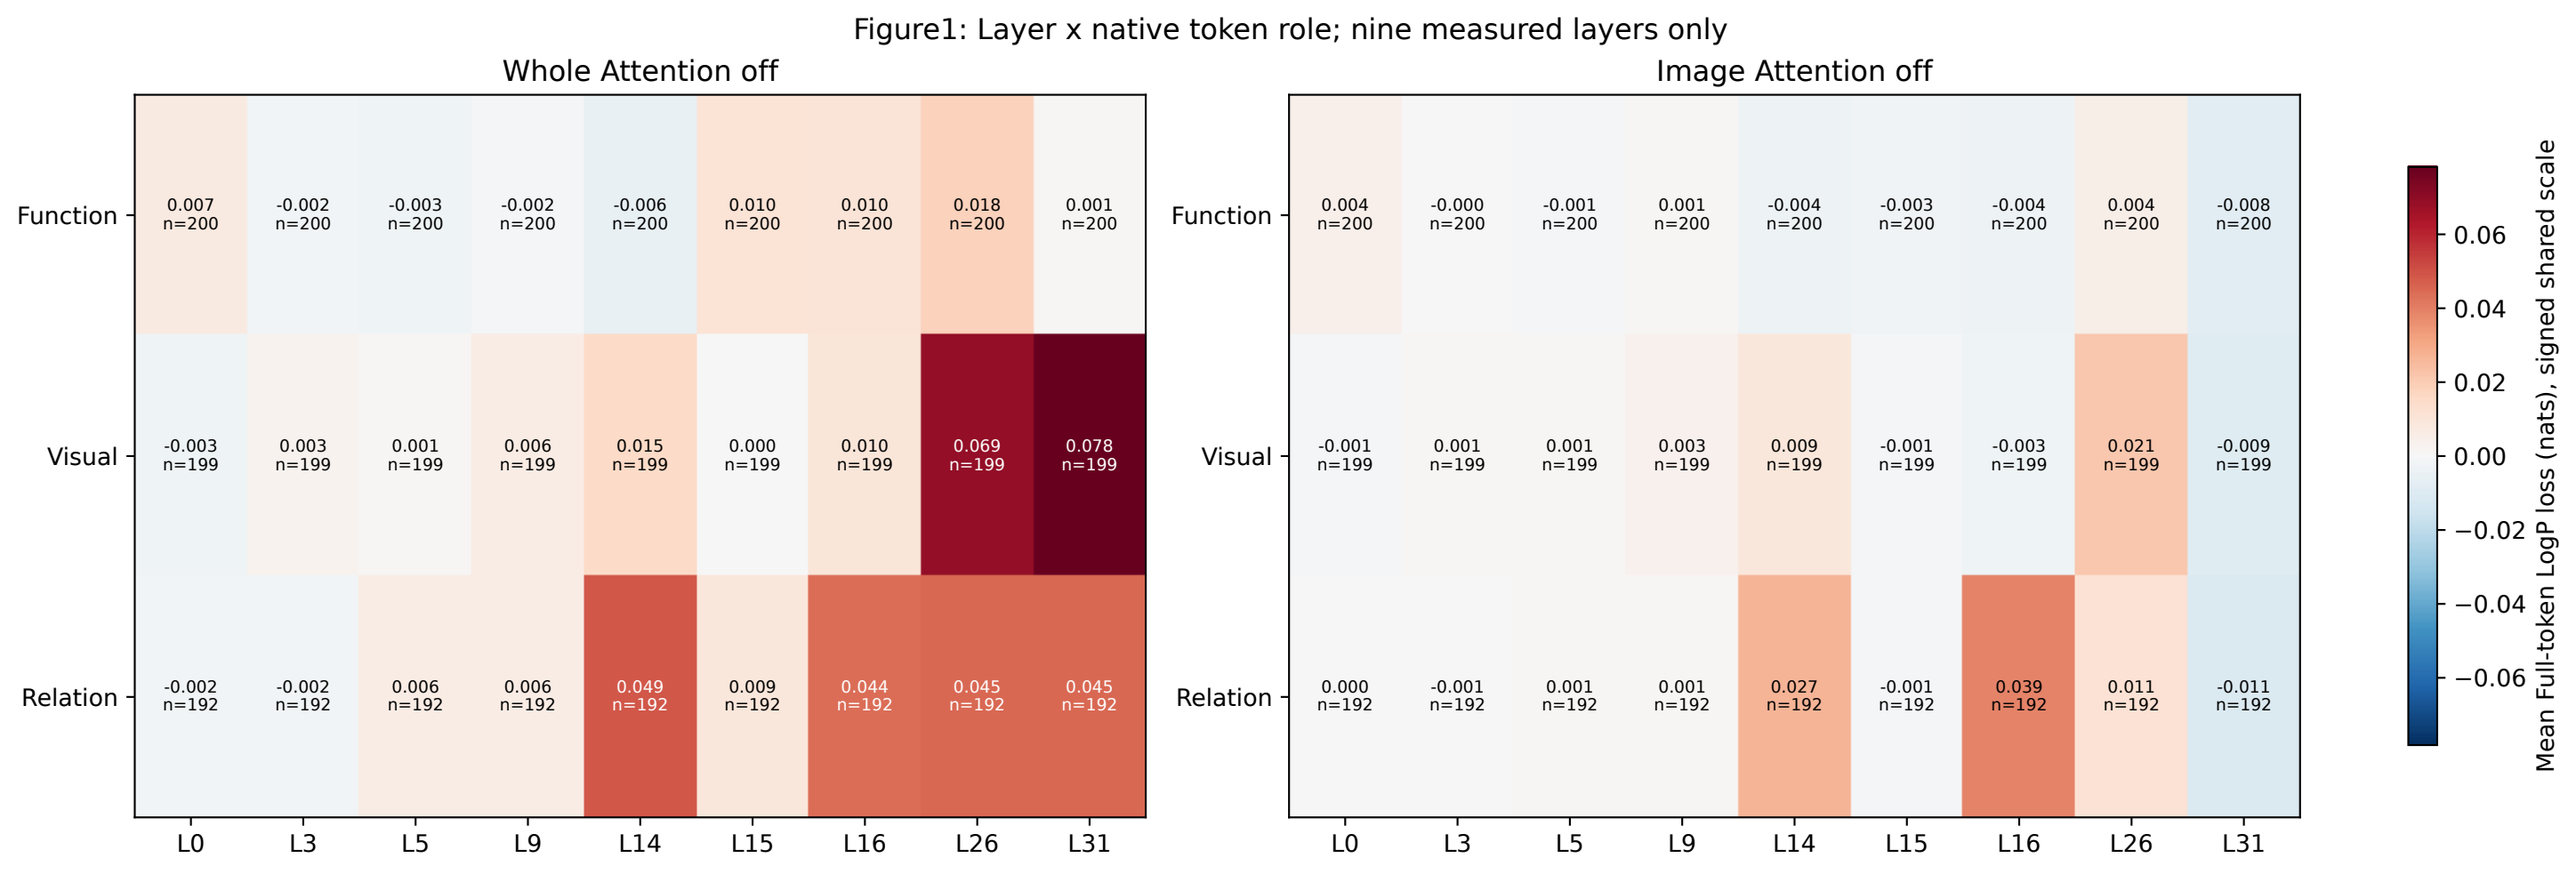
Figure1: Layer x native token role; nine measured layers only
Whole Attention off
Image Attention off
Function
Visual
Relation
L0
L3
L5
L9
L14
L15
L16
L26
L31
0.007 n=200
-0.002 n=200
-0.003 n=200
-0.002 n=200
-0.006 n=200
0.010 n=200
0.010 n=200
0.018 n=200
0.001 n=200
-0.003 n=199
0.003 n=199
0.001 n=199
0.006 n=199
0.015 n=199
0.000 n=199
0.010 n=199
0.069 n=199
0.078 n=199
-0.002 n=192
-0.002 n=192
0.006 n=192
0.006 n=192
0.049 n=192
0.009 n=192
0.044 n=192
0.045 n=192
0.045 n=192
0.004 n=200
-0.000 n=200
-0.001 n=200
0.001 n=200
-0.004 n=200
-0.003 n=200
-0.004 n=200
0.004 n=200
-0.008 n=200
-0.001 n=199
0.001 n=199
0.001 n=199
0.003 n=199
0.009 n=199
-0.001 n=199
-0.003 n=199
0.021 n=199
-0.009 n=199
0.000 n=192
-0.001 n=192
0.001 n=192
0.001 n=192
0.027 n=192
-0.001 n=192
0.039 n=192
0.011 n=192
-0.011 n=192
0.06
0.04
0.02
0.00
−0.02
−0.04
−0.06
Mean Full-token LogP loss (nats), signed shared scale