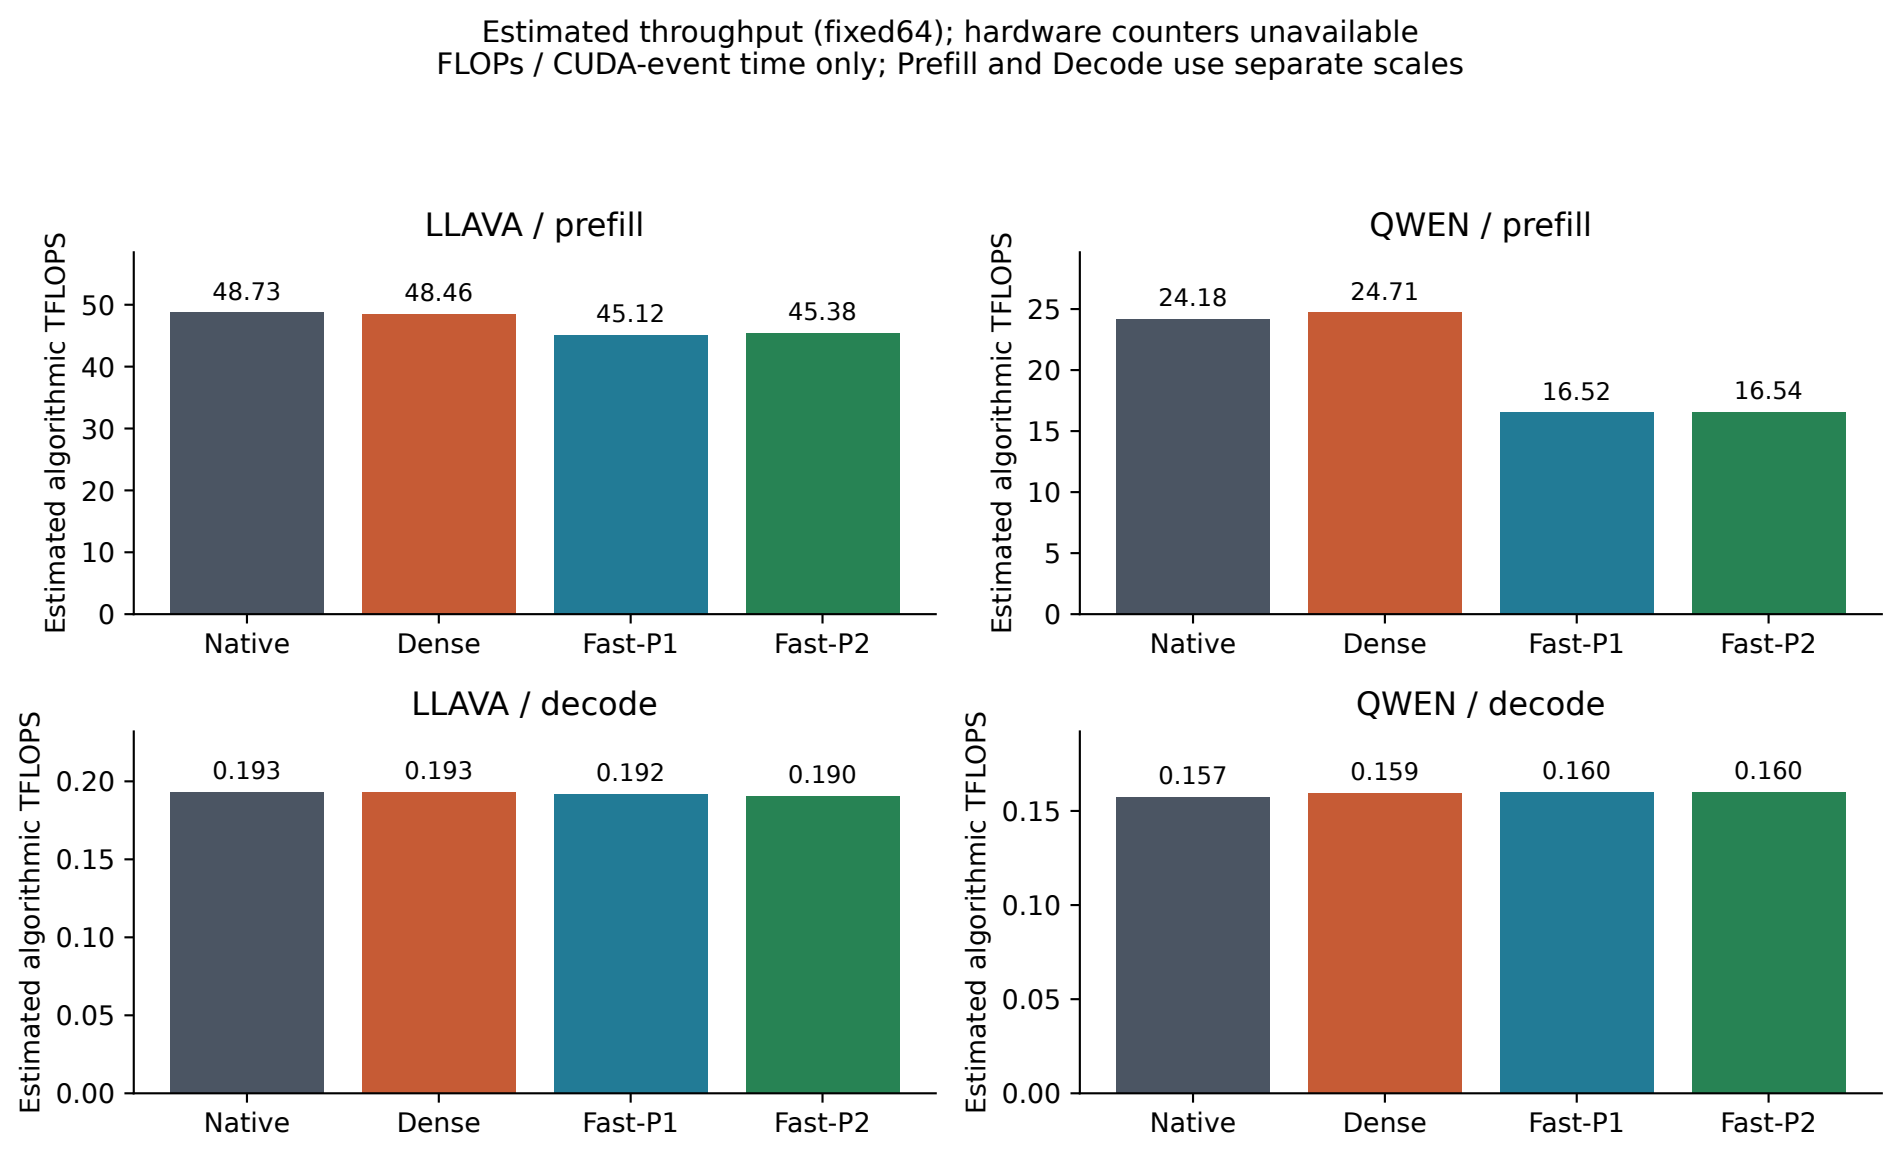
Estimated throughput (fixed64); hardware counters unavailable
FLOPs / CUDA-event time only; Prefill and Decode use separate scales
LLAVA / prefill
QWEN / prefill
LLAVA / decode
QWEN / decode
Estimated algorithmic TFLOPS
48.73
48.46
45.12
45.38
24.18
24.71
16.52
16.54
0.193
0.193
0.192
0.190
0.157
0.159
0.160
0.160
Native
Dense
Fast-P1
Fast-P2