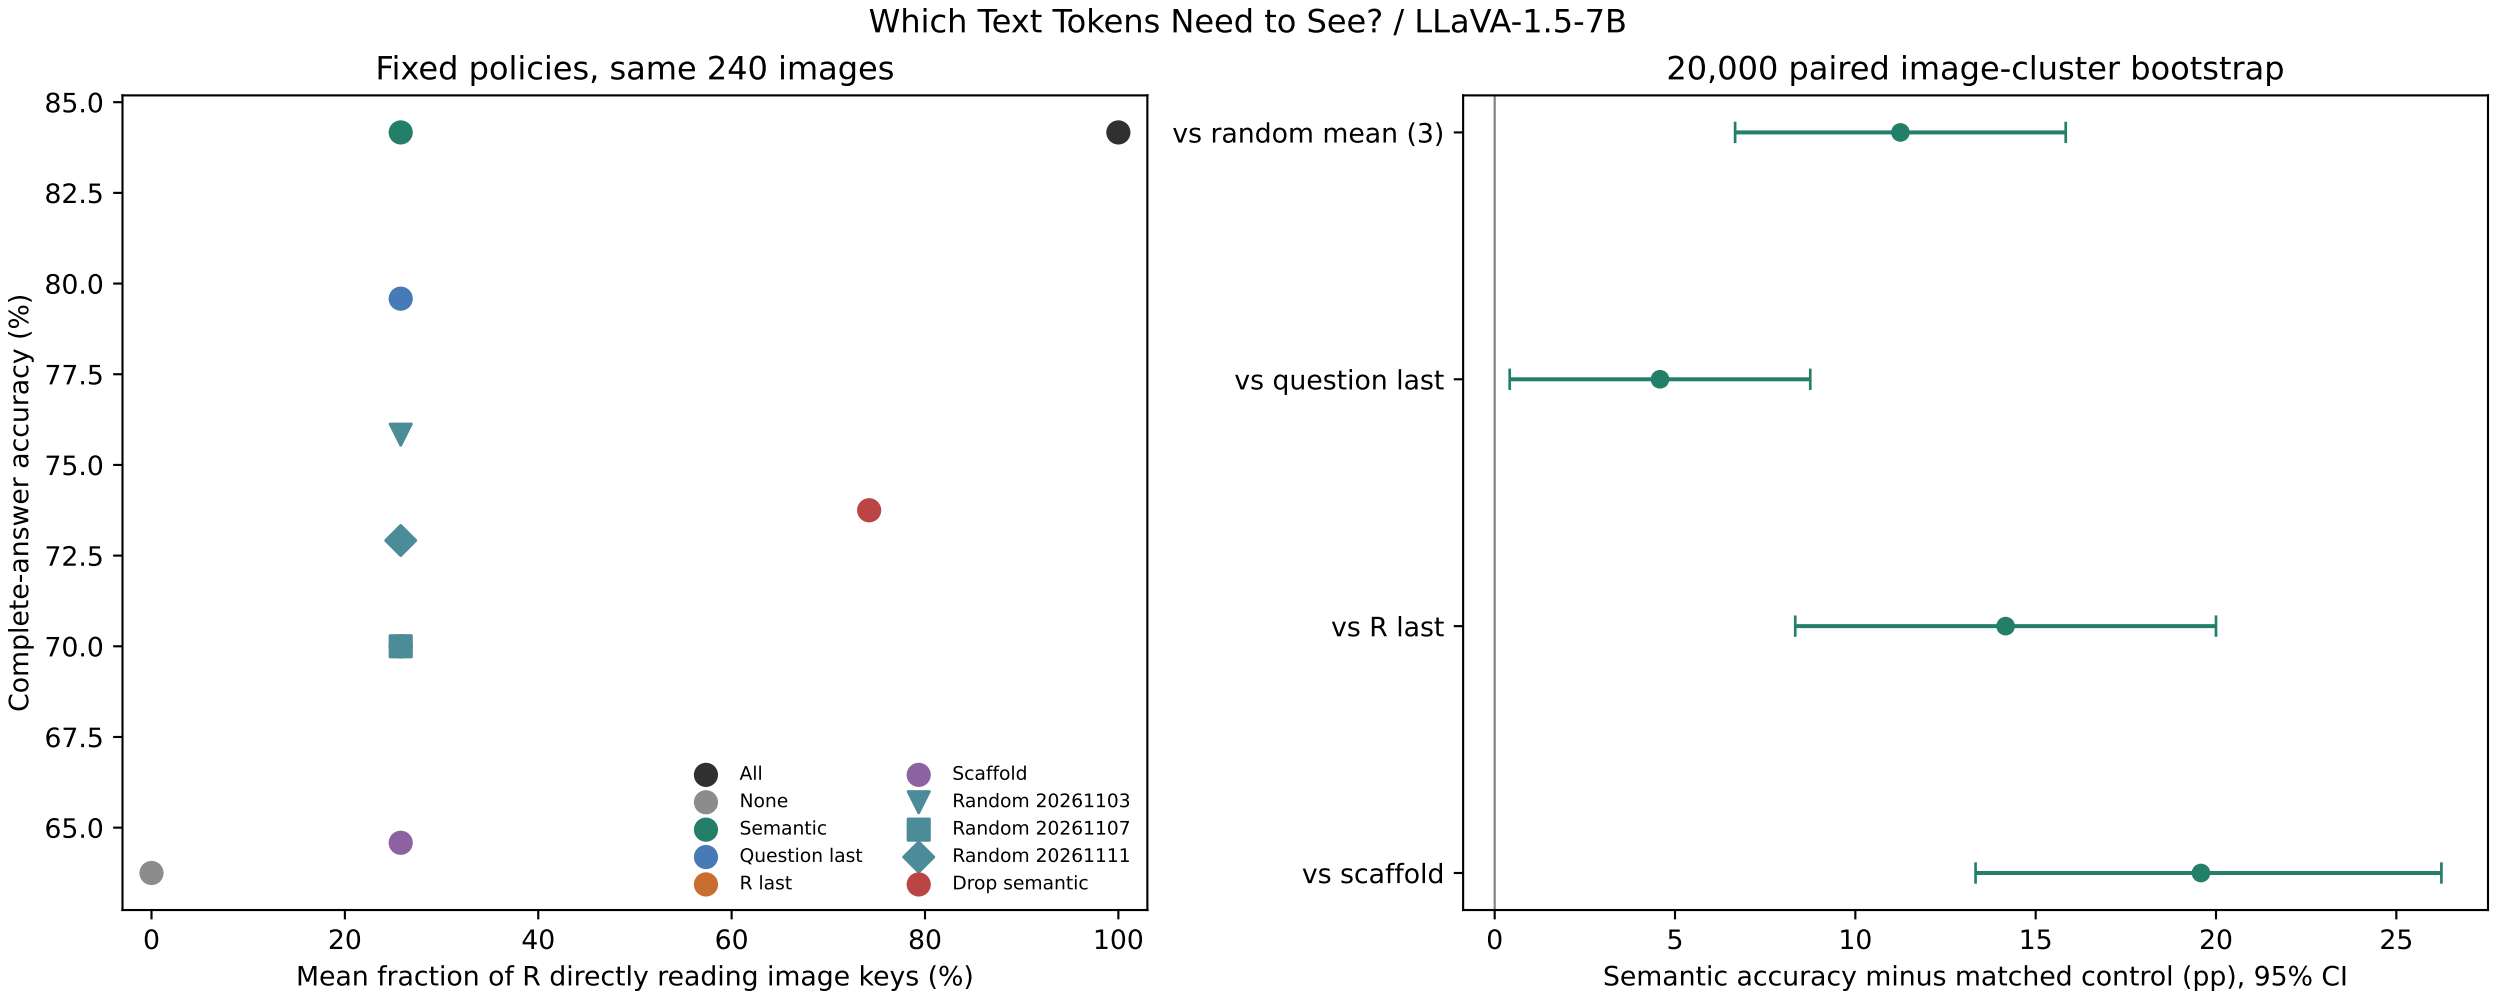
Which Text Tokens Need to See? / LLaVA-1.5-7B
Fixed policies, same 240 images
Complete-answer accuracy (%)
85.0
82.5
80.0
77.5
75.0
72.5
70.0
67.5
65.0
0
20
40
60
80
100
Mean fraction of R directly reading image keys (%)
All
None
Semantic
Question last
R last
Scaffold
Random 20261103
Random 20261107
Random 20261111
Drop semantic
20,000 paired image-cluster bootstrap
vs random mean (3)
vs question last
vs R last
vs scaffold
0
5
10
15
20
25
Semantic accuracy minus matched control (pp), 95% CI