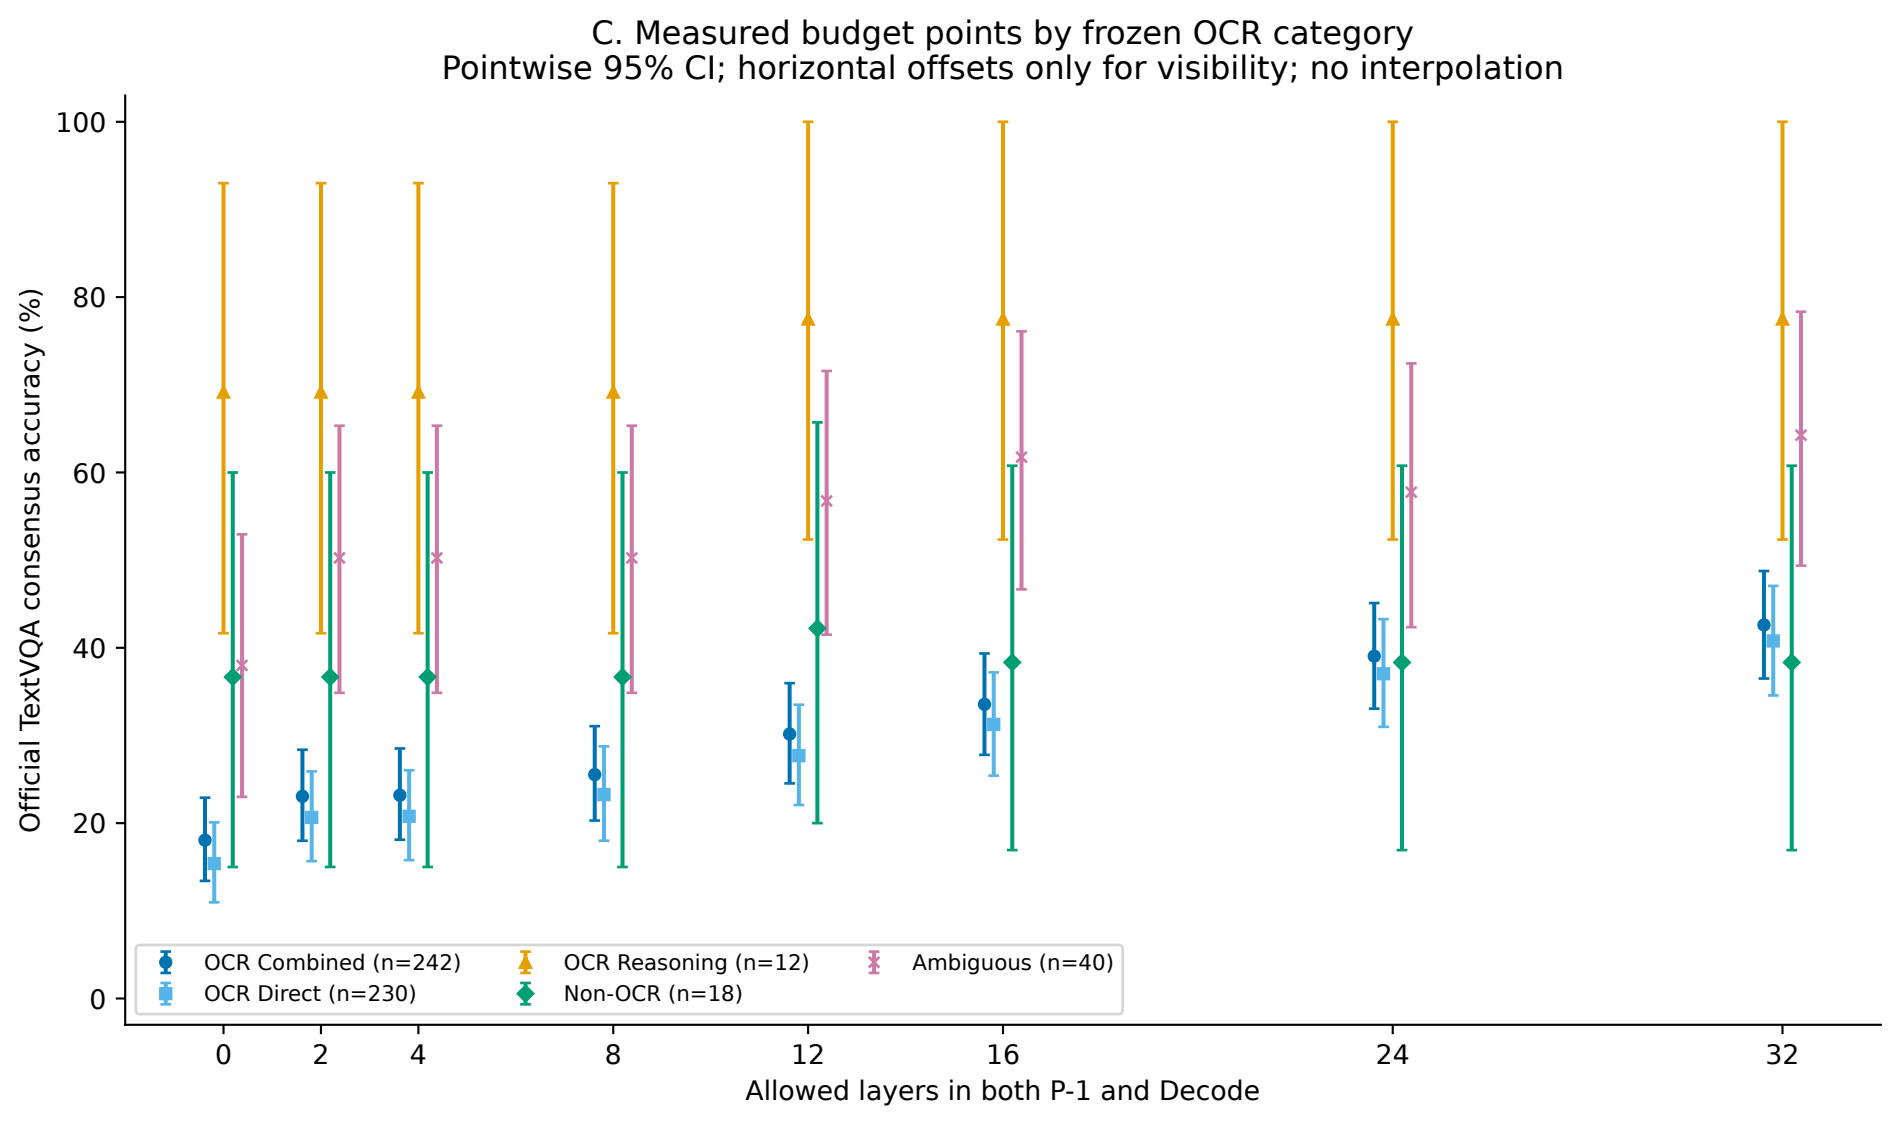
C. Measured budget points by frozen OCR category
Pointwise 95% CI; horizontal offsets only for visibility; no interpolation
Official TextVQA consensus accuracy (%)
100
80
60
40
20
0
0
2
4
8
12
16
24
32
Allowed layers in both P-1 and Decode
OCR Combined (n=242)
OCR Direct (n=230)
OCR Reasoning (n=12)
Non-OCR (n=18)
Ambiguous (n=40)